

# Problem Solving with AI Techniques

## Stochastic Search

Paul Weng

UM-SJTU Joint Institute

VE593, Fall 2018



JOINT INSTITUTE  
交大密西根学院

## 1 Monte Carlo Tree Search

- Monte Carlo Methods
- Principle of Monte Carlo Tree Search

## 2 Simulated Annealing

- Local Search
- Principle of Simulated Annealing

# Monte Carlo Methods

- Monte Carlo method = sampling from a probability distribution
- It allows to estimate an integral (e.g.,  $\mathbb{E}$  is an integral)  
e.g., If  $\mathbb{E}_p[f(X)] = \int_x f(x)p(x)dx$ , then

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_i f(x_i) = \mathbb{E}_p[f(X)]$$

# Monte Carlo Methods

- Monte Carlo method = sampling from a probability distribution
- It allows to estimate an integral (e.g.,  $\mathbb{E}$  is an integral)  
e.g., If  $\mathbb{E}_p[f(X)] = \int_x f(x)p(x)dx$ , then

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_i f(x_i) = \mathbb{E}_p[f(X)]$$

or

$$\lim_{N \rightarrow \infty} \sum_i w_i f(x_i) = \mathbb{E}_p[f(X)]$$

- Useful when  $p$  is complicated and  $\mathbb{E}_p[f(X)]$  cannot be calculated directly

# Rejection Sampling

- How can we generate i.i.d. samples  $x_i \sim p(x)$ ?
- **Assumptions:**
  - We can sample  $x \sim q(x)$  from a simpler distribution  $q(x)$  (e.g., uniform), called **proposal distribution**
  - We can numerically evaluate  $p(x)$  for specific  $x$  (even if we don't have an analytic expression of  $p(x)$ )
  - There exists  $M$  such that  $\forall x, p(x) \leq Mq(x)$  (which implies  $q$  has a larger or equal support as  $p$ )
- **Rejection Sampling:**
  - Sample a candidate  $x \sim q(x)$
  - Accept  $x$  with probability  $\frac{p(x)}{Mq(x)}$  and reject otherwise
  - Repeat until sample size =  $N$
- This generates an unweighted sample set to approximate  $p(x)$

# Importance Sampling

- **Assumptions:**

- We can sample  $x \sim q(x)$  from a simpler distribution  $q(x)$  (e.g., uniform)
- We can numerically evaluate  $p(x)$  for specific  $x$  (even if we don't have an analytic expression of  $p(x)$ )

- **Importance Sampling:**

- Sample a candidate  $x_i \sim q(x)$
- Add the weighted sample  $(w_i, x_i)$  where  $w_i = \frac{p(x_i)}{q(x_i)}$
- Repeat  $N$  times
- This generates a weighted sample set to approximate  $p(x)$   
Weights  $w_i$  are called importance weights
- **Crucial for efficiency:** a good choice of proposal distribution  $q(x)$

- 1 Monte Carlo Tree Search
  - Monte Carlo Methods
  - Principle of Monte Carlo Tree Search
  
- 2 Simulated Annealing
  - Local Search
  - Principle of Simulated Annealing

# Monte Carlo Tree Search (MTCS)

- MCTS is very successful on Computer Go and other games
- MCTS is one of the main components in Alphago
- MCTS is a quite novel technique ( $\sim 10$  years)
- MCTS is rather simple to implement
- MCTS is very general: applicable on any discrete domain



# Flat Monte Carlo


- Goal of MCTS: estimate the expected value of action (*Q-function*)

$$Q(s, a) = \mathbb{E}[\Delta \mid s, a]$$

where expectation is taken over future randomized actions (including possible adversary choices and possible stochastic transitions in the environment)

- In search trees, MCTS provides an estimate of  $f$
- *Flat Monte Carlo* does so by rolling out many random simulations (using a `ROLLOUTPOLICY`) without growing a tree
- Key difference/advantage of MCTS over flat Monte Carlo: search focuses computational effort on promising actions

# Generic MCTS Scheme



from Browne et al.

- 1: start tree  $V = \{v_0\}$
- 2: **while** within computational budget **do**
- 3:    $v_l \leftarrow \text{TREEPOLICY}(V)$  chooses a leaf of  $V$
- 4:   append  $v_l$  to  $V$
- 5:    $\Delta \leftarrow \text{ROLLOUTPOLICY}(V)$  rolls out a full simulation, with return  $\Delta$
- 6:    $\text{BACKUP}(v_l, \Delta)$  updates the values of all parents of  $v_l$
- 7: **end while**
- 8: return best child of  $v_0$

# Generic MCTS Scheme: Remarks

- Like flat MC, MCTS typically computes full roll-outs to a terminal state. A heuristic to estimate the utility of a state is not needed, but can be incorporated.
- The tree grows unbalanced.
- TREEPOLICY decides where the tree is expanded – and needs to trade-off exploration vs. exploitation.
- ROLLOUTPOLICY is necessary to simulate a roll-out. It typically is a random policy (at least a randomized policy).

# Upper Confidence Tree (UCT)

- UCT uses UCB to realize TREEPOLICY, i.e., to decide where to expand the tree
- BACKUP updates all parents of  $v_l$  as
$$n(v) \leftarrow n(v) + 1 \text{ (count how often it has been tried)}$$
$$Q(v) \leftarrow Q(v) + \Delta \text{ (sum of rewards received)}$$
- TREEPOLICY chooses child nodes based on UCB:

$$\arg \max \frac{Q(v')}{n(v')} + \beta \sqrt{\frac{2 \ln n(v)}{n(v')}}}$$

or chooses  $v'$  if  $n(v') = 0$

## 1 Monte Carlo Tree Search

- Monte Carlo Methods
- Principle of Monte Carlo Tree Search

## 2 Simulated Annealing

- Local Search
- Principle of Simulated Annealing


# Iterative Improvement Algorithms

- In some problems, *path* is irrelevant, *goal state* is the solution

$$\begin{array}{rcccccc}
 & & S & E & N & D \\
 + & & M & O & R & E \\
 \hline
 = & M & O & N & E & Y
 \end{array}$$


- State space = set of *complete* configurations  
 find configuration satisfying all constraints, e.g., timetable  
 find *optimal* configuration, e.g., Travelling Salesperson Problem
- Iterative improvement algorithms:** keep a single "current" state, try to improve it
- Constant space complexity

# Example: $n$ -Queens



- **Goal:** Place  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal
- **Algorithmic principle:**
  - Start with  $n$  queens placed on board
  - Repeatedly, move a queen to reduce number of conflicts


# Example: $n$ -Queens



- **Goal:** Place  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal
- **Algorithmic principle:**
  - Start with  $n$  queens placed on board
  - Repeatedly, move a queen to reduce number of conflicts
- Can solve  $n$ -queens problems almost instantaneously for very large  $n$ , e.g.,  $n = 10^6$




# Example: Travelling Salesperson Problem



- **Goal:** Find minimal-cost cycle
- **Algorithmic principle:**
  - Start with any complete tour
  - Repeatedly, perform pairwise exchanges

# Example: Travelling Salesperson Problem



- **Goal:** Find minimal-cost cycle
- **Algorithmic principle:**
  - Start with any complete tour
  - Repeatedly, perform pairwise exchanges
- Variants of this approach get within 1% of optimal very quickly with thousands of cities

# Hill-climbing (or gradient ascent/descent)


”Like climbing Everest in thick fog with amnesia”  
(assuming we maximize an utility function)

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                   neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
  end
```


# Hill-climbing (contd.)

- Hill-climbing can get stuck, see **state space landscape**



# Hill-climbing (contd.)


- Hill-climbing can get stuck, see **state space landscape**



- Random-restart hill climbing:**  
overcomes local maxima — trivially complete

# Hill-climbing (contd.)

- Hill-climbing can get stuck, see **state space landscape**



- Random-restart hill climbing:**  
overcomes local maxima — trivially complete
- Random sideways moves:**  
 $\oplus$  escape from shoulders,  $\ominus$  loop on flat maxima

- 1 Monte Carlo Tree Search
  - Monte Carlo Methods
  - Principle of Monte Carlo Tree Search
- 2 Simulated Annealing
  - Local Search
  - Principle of Simulated Annealing

# Simulated Annealing

- **Idea:** Escape local maxima by allowing some “bad” moves  
*but gradually decrease their size and frequency*

**function** **SIMULATED-ANNEALING**( *problem*, *schedule* ) **returns** a solution state

**inputs:** *problem*, a problem

*schedule*, a mapping from time to “temperature”

**local variables:** *current*, a node

*next*, a node

*T*, a “temperature” controlling prob. of downward steps

*current*  $\leftarrow$  MAKE-NODE(INITIAL-STATE[*problem*])

**for** *t*  $\leftarrow$  1 **to**  $\infty$  **do**

*T*  $\leftarrow$  *schedule*[*t*]

**if** *T* = 0 **then return** *current*

*next*  $\leftarrow$  a randomly selected successor of *current*

$\Delta E \leftarrow$  VALUE[*next*] – VALUE[*current*]

**if**  $\Delta E > 0$  **then** *current*  $\leftarrow$  *next*

**else** *current*  $\leftarrow$  *next* only with probability  $e^{\Delta E / T}$



# Properties of simulated annealing

- At fixed “temperature”  $T$ , state occupation probability reaches Boltzman distribution

$$p(x) \propto e^{\frac{E(x)}{kT}}$$

$T$  decreased slowly enough  $\implies$  always reach best state  $x^*$

because  $e^{\frac{E(x^*)}{kT}} / e^{\frac{E(x)}{kT}} = e^{\frac{E(x^*) - E(x)}{kT}} \gg 1$  for small  $T$

- Is this necessarily an interesting guarantee?
- Devised by Metropolis et al., 1953, for physical process modelling
- Widely used in VLSI layout, airline scheduling, etc.

# Going Further...

- **Local beam search:** keep  $k$  states instead of 1; choose top  $k$  of all their successors
- **Genetic algorithms:** stochastic local beam search + generate successors from *pairs* of states
- **Evolutionary Strategies** (e.g., CMA-ES)
- **Nature-inspired algorithms** (e.g., Particle Swarm Optimization)
- **Metaheuristics** (e.g., Cross-Entropy Method)