

# Problem Solving with AI Techniques

## Artificial Neural Networks

Paul Weng

UM-SJTU Joint Institute

VE593, Fall 2018



JOINT INSTITUTE  
交大密西根学院

# 1 Introduction

- Definitions
- Why ANN?



## 2 Inference

## 3 Learning

## 4 Discussion

- Architecture Design
- How to Obtain Better Performances with ANNs?

# Artificial Neuron



- **Perceptron:**  $f(x) = \text{sign}(x)$
- **Logistic regression:**  $f(x) = \frac{1}{1+e^{-x}}$
- **Linear regression:**  $f(x) = x$
- $f$  called **activation function**, usually it is (sub)differentiable
- Why non-linear activation function?

# Examples of Activation Functions

- logistic or sigmoid function

- $f(x) = \frac{1}{1+e^{-x}}$
- $f'(x) = f(x)(1 - f(x))$

- tanh




- $f(x) = \frac{1-e^{-2x}}{1+e^{-2x}}$
- $f'(x) = 1 - f(x)^2$

- ReLU (Rectified Linear Unit)

- $f(x) = \max(0, x)$
- $f'(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{otherwise} \end{cases}$

- leaky ReLU

- $f(x) = \max(-\beta x, x)$  with small  $\beta > 0$
- $f'(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ -\beta & \text{otherwise} \end{cases}$




from wikipedia

# Artificial Neural Networks (ANN)

- **ANN** is defined as a directed weighted graph:
  - three types of nodes: input nodes (nodes without predecessors), output nodes (nodes without successors) and hidden nodes
  - an activation function is associated to each non-input node
  - a weight is associated to each edge
  - **Computation:**
    - Input nodes send a fixed value to their successors
    - Non-input nodes compute weighted sum of their inputs and send results transformed by their activation function to their successors
- Feedforward vs recurrent ANNs
- **ANN with fixed architecture** defines a class of functions parametrized by the ANN's weights
- **ANN** is a composition of functions
- **Multi-Layer Perceptron:** graph is organized in layers

# Multi-Layer Perceptron (MLP)

- **Multi-Layer Perceptron** = ANN whose graph is organized in layers + any non-input node connected to all nodes of its previous layer



from Amazon

## 1 Introduction

- Definitions
- Why ANN?

## 2 Inference

## 3 Learning

## 4 Discussion

- Architecture Design
- How to Obtain Better Performances with ANNs?


# Empirical Argument

- State-of-the-art performance in some domains:
  - Computer vision
  - Speech processing
  - Natural language processing
- Possible thanks to:
  - Very large dataset
  - Powerful hardware
  - Efficient techniques for training ANNs (and especially deep learning)



# CS Theoretical Argument

- Perceptrons can implement a NAND gate:



- Any logic function can be performed using only NAND gates
- Therefore ANN can perform any computation


# Mathematical Theoretical Argument

- ANN is a universal function approximator:

$\forall f : \mathbb{R}^D \rightarrow \mathbb{R}, \forall \varepsilon > 0, \exists$  an MLP with one hidden layer  $g : \mathbb{R}^D \rightarrow \mathbb{R}$ ,

$$\forall \mathbf{x} \in \mathbb{R}^D, |f(\mathbf{x}) - g(\mathbf{x})| \leq \varepsilon$$

- Illustration:** ANN can approximate a step function



where  $z_i = w_i x + b_i = w_i(x + b_i/w_i)$  with  $w_i$  large

- 1 Introduction
  - Definitions
  - Why ANN?


- 2 Inference

- 3 Learning

- 4 Discussion
  - Architecture Design
  - How to Obtain Better Performances with ANNs?

# Notations for MLP

- $n^l = \#$  of nodes at layer  $l = 1, \dots, L$
- $\mathbf{a}^l$  = vector of outputs of the  $l$ -th layer of an MLP. Therefore,  $\mathbf{a}^1 = \mathbf{x}$
- $w_{ij}^l$  = weight of connection between node  $i$  of layer  $l$  and node  $j$  of layer  $l - 1$
- Focusing on the  $i$ -th node of layer  $l = 2, \dots, L$  of an MLP



- In matrix notations,  $\mathbf{a}^l = f^l(\mathbf{z}^l) = f^l(\mathbf{w}^l \mathbf{a}^{l-1} + \mathbf{b}^l) = g^l(\mathbf{a}^{l-1})$

# Forward Pass: Illustration

- Therefore  $\mathbf{a}^L = g^L(g^{L-1}(\dots g^2(\mathbf{a}^1)))$
- Example on blackboard

# Forward Pass: Algorithm

- Assuming

$$\mathbf{W} = (\mathbf{w}^2, \dots \mathbf{w}^L)$$

$$\mathbf{B} = (\mathbf{b}^2, \dots \mathbf{b}^L)$$

$$\mathbf{F} = (\mathbf{f}^2, \dots \mathbf{f}^L)$$

where  $\mathbf{w}^l \in \mathbb{R}^{n_l \times n_{l-1}}$ ,  $\mathbf{b}^l \in \mathbb{R}^{n_l}$ , and  $f^l : \mathbb{R} \rightarrow \mathbb{R}$

---

```

1 Forward( $\mathbf{x}, \mathbf{W}, \mathbf{B}, \mathbf{F}$ )
2  $\mathbf{a}^0 \leftarrow \mathbf{x}$ 
3 for  $l = 2$  to  $L$  do
4    $\mathbf{z}^l \leftarrow \mathbf{w}^l \mathbf{a}^{l-1} + \mathbf{b}^l$ 
5    $\mathbf{a}^l \leftarrow \mathbf{f}^l(\mathbf{z}^l)$ 
6 return  $\mathbf{a}^L$ 
  
```

---

# 1 Introduction

- Definitions
- Why ANN?

## 2 Inference

## 3 Learning

## 4 Discussion

- Architecture Design
- How to Obtain Better Performances with ANNs?

# General Problem

- Recall ERM (possibly with regularization)

$$R_{\mathcal{D}}(\mathbf{W}, \mathbf{B}) - \rho(\mathbf{W}, \mathbf{B}) = \frac{1}{N} \sum_{n=1}^N \ell(g_{\mathbf{W}, \mathbf{B}}(\mathbf{x}^n), y^n) - \rho(\mathbf{W}, \mathbf{B})$$

where  $g_{\mathbf{W}, \mathbf{B}}$  is the function computed by ANN with parameters  $\mathbf{W}, \mathbf{B}$



# General Problem

- Recall ERM (possibly with regularization)

$$R_{\mathcal{D}}(\mathbf{W}, \mathbf{B}) - \rho(\mathbf{W}, \mathbf{B}) = \frac{1}{N} \sum_{n=1}^N \ell(g_{\mathbf{W}, \mathbf{B}}(\mathbf{x}^n), y^n) - \rho(\mathbf{W}, \mathbf{B})$$

where  $g_{\mathbf{W}, \mathbf{B}}$  is the function computed by ANN with parameters  $\mathbf{W}, \mathbf{B}$

- Method: (Mini-batch stochastic) gradient descent

# General Problem

- Recall ERM (possibly with regularization)

$$R_{\mathcal{D}}(\mathbf{W}, \mathbf{B}) - \rho(\mathbf{W}, \mathbf{B}) = \frac{1}{N} \sum_{n=1}^N \ell(g_{\mathbf{W}, \mathbf{B}}(\mathbf{x}^n), y^n) - \rho(\mathbf{W}, \mathbf{B})$$

where  $g_{\mathbf{W}, \mathbf{B}}$  is the function computed by ANN with parameters  $\mathbf{W}, \mathbf{B}$

- Method: (Mini-batch stochastic) gradient descent
- Issues:
  - Non-convex optimization problem
  - How to compute the gradient?

# Backpropagation: Introduction

- **ANN** is a composition of functions  $g_{\mathbf{W}, \mathbf{B}} = g^L \circ g^{L-1} \circ \dots \circ g^2$
- Gradient can be computed by the **chain rule**
- Closed-form equation of gradient for an ANN may be complex, but
- Its value at a fixed  $(\mathbf{W}, \mathbf{B})$  can be computed recursively!
- This recursive computation is called **backpropagation**

# Backpropagation: Last Layer $L$

- For node  $j$  of layer  $L$ ,

$$\begin{aligned}
 \frac{\partial \ell(\mathbf{a}^L, \mathbf{y})}{\partial w_{jk}^L} &= \frac{\partial \ell(f^L(\mathbf{z}^L), \mathbf{y})}{\partial w_{jk}^L} \\
 &= \frac{\partial z_j^L}{\partial w_{jk}^L} f^{L'}(z_j^L) \frac{\partial \ell(\mathbf{a}^L, \mathbf{y})}{\partial a_j} \\
 &= a_k^{L-1} f^{L'}(z_j^L) \frac{\partial \ell(\mathbf{a}^L, \mathbf{y})}{\partial a_j}
 \end{aligned}$$

$$\frac{\partial R}{\partial w_{jk}^L} = a_k^{L-1} f^{L'}(z_j^L) \frac{\partial R}{\partial a_j}$$

$$\frac{\partial R}{\partial b_j^L} = f^{L'}(z_j^L) \frac{\partial R}{\partial a_j}$$

- Let  $\delta^L = f^{L'}(\mathbf{z}^L) \otimes \nabla_{\mathbf{a}} R = \frac{\partial R}{\partial \mathbf{z}^L}$

# Backpropagation: Recurrence

- At node  $j$  of layer  $l$ ,

$$\delta_j^l = \frac{\partial R}{\partial z_j^l} = \sum_k \frac{\partial R}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} = \sum_k \frac{\partial z_k^{l+1}}{\partial z_j^l} \delta_k^{l+1}$$

# Backpropagation: Recurrence

- At node  $j$  of layer  $l$ ,

$$\delta_j^l = \frac{\partial R}{\partial z_j^l} = \sum_k \frac{\partial R}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} = \sum_k \frac{\partial z_k^{l+1}}{\partial z_j^l} \delta_k^{l+1}$$

- By definition,  $z_k^{l+1} = \sum_j w_{kj}^{l+1} a_j^l + b_k^{l+1} = \sum_j w_{kj}^{l+1} f^l(z_j^l) + b_k^{l+1}$

# Backpropagation: Recurrence

- At node  $j$  of layer  $l$ ,

$$\delta_j^l = \frac{\partial R}{\partial z_j^l} = \sum_k \frac{\partial R}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} = \sum_k \frac{\partial z_k^{l+1}}{\partial z_j^l} \delta_k^{l+1}$$

- By definition,  $z_k^{l+1} = \sum_j w_{kj}^{l+1} a_j^l + b_k^{l+1} = \sum_j w_{kj}^{l+1} f^l(z_j^l) + b_k^{l+1}$
- By differentiating,  $\frac{\partial z_k^{l+1}}{\partial z_j^l} = w_{kj}^{l+1} f'^l(z_j^l)$

# Backpropagation: Recurrence

- At node  $j$  of layer  $l$ ,

$$\delta_j^l = \frac{\partial R}{\partial z_j^l} = \sum_k \frac{\partial R}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} = \sum_k \frac{\partial z_k^{l+1}}{\partial z_j^l} \delta_k^{l+1}$$

- By definition,  $z_k^{l+1} = \sum_j w_{kj}^{l+1} a_j^l + b_k^{l+1} = \sum_j w_{kj}^{l+1} f'(z_j^l) + b_k^{l+1}$
- By differentiating,  $\frac{\partial z_k^{l+1}}{\partial z_j^l} = w_{kj}^{l+1} f''(z_j^l)$
- Therefore,

$$\delta_j^l = \sum_k w_{kj}^{l+1} f''(z_j^l) \delta_k^{l+1} = \mathbf{w}_{\cdot j}^{l+1 \top} \boldsymbol{\delta}^{l+1} f''(z_j^l)$$

$$\boldsymbol{\delta}^l = \mathbf{w}^{l+1 \top} \boldsymbol{\delta}^{l+1} \otimes f''(\mathbf{z}^l)$$



# Backpropagation: Layer $l$

- At node  $j$  of layer  $l$ ,  $\frac{\partial R}{\partial w_{jk}^l} = \frac{\partial R}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l}$

# Backpropagation: Layer $l$

- At node  $j$  of layer  $l$ ,  $\frac{\partial R}{\partial w_{jk}^l} = \frac{\partial R}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l}$
- By definition,  $z_j^l = \sum_k w_{jk}^l a_k^{l-1} + b_j^l$

# Backpropagation: Layer $l$

- At node  $j$  of layer  $l$ ,  $\frac{\partial R}{\partial w_{jk}^l} = \frac{\partial R}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l}$
- By definition,  $z_j^l = \sum_k w_{jk}^l a_k^{l-1} + b_j^l$
- By differentiating,  $\frac{\partial z_j^l}{\partial w_{jk}^l} = a_k^{l-1}$

# Backpropagation: Layer $l$

- At node  $j$  of layer  $l$ ,  $\frac{\partial R}{\partial w_{jk}^l} = \frac{\partial R}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l}$
- By definition,  $z_j^l = \sum_k w_{jk}^l a_k^{l-1} + b_j^l$
- By differentiating,  $\frac{\partial z_j^l}{\partial w_{jk}^l} = a_k^{l-1}$
- Therefore,  $\frac{\partial R}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$

# Backpropagation: Layer $l$

- At node  $j$  of layer  $l$ ,  $\frac{\partial R}{\partial w_{jk}^l} = \frac{\partial R}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l}$
- By definition,  $z_j^l = \sum_k w_{jk}^l a_k^{l-1} + b_j^l$
- By differentiating,  $\frac{\partial z_j^l}{\partial w_{jk}^l} = a_k^{l-1}$
- Therefore,  $\frac{\partial R}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$
- Similarly,  $\frac{\partial R}{\partial b_j^l} = \delta_j^l$

# Backpropagation: Summary

The gradient of the ANN at current parameter can be computed with:

- $\delta_j^L = \frac{\partial R}{\partial a_j^L} f''(z_j^L)$
- $\delta^l = (\mathbf{w}^{l+1\top} \delta^{l+1}) \otimes f''(\mathbf{z}^l)$
- $\frac{\partial R}{\partial b_j^l} = \delta_j^l$
- $\frac{\partial R}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$

# Backpropagation: Algorithm

---



---

```

1 Backpropagation( $\mathbf{x}, \mathbf{W}, \mathbf{B}, \mathbf{F}$ )
2  $\mathbf{a}^1 \leftarrow \mathbf{x}$ 
3 for  $l = 2$  to  $L$  do
4    $\mathbf{z}^l \leftarrow \mathbf{w}^l \mathbf{a}^{l-1} + \mathbf{b}^l$ 
5    $\mathbf{a}^l \leftarrow f^l(\mathbf{z}^l)$ 
6  $\delta^L \leftarrow \nabla_a R \otimes f^{L'}(\mathbf{z}^L)$ 
7 for  $l = L - 1$  to  $2$  do
8    $\delta^l \leftarrow \mathbf{w}^{l+1\top} \delta^{l+1} \otimes f^{l'}(\mathbf{z}^l)$ 
9 return  $\frac{\partial R}{\partial \mathbf{w}_{jk}^l} = \mathbf{a}_k^{l-1} \delta_j^l$  and  $\frac{\partial R}{\partial \mathbf{b}_j^l} = \delta_j^l$ 
  
```

---

## 1 Introduction

- Definitions
- Why ANN?

## 2 Inference

## 3 Learning

## 4 Discussion

- Architecture Design
- How to Obtain Better Performances with ANNs?



# How to Choose the Number of Layers/Neurons?

Two schools:

- Start by training small network
- Grow it by increasing # nodes in a layer or adding layers
- Stop when performance on test set doesn't improve anymore

or

- Start with as large network as you can
- Train with regularized ERM

## 1 Introduction

- Definitions
- Why ANN?

## 2 Inference

## 3 Learning

## 4 Discussion

- Architecture Design
- How to Obtain Better Performances with ANNs?

# Regularization

- Recall

$$\min_{\mathbf{W}, \mathbf{B}} R_{\mathcal{D}}(\mathbf{W}, \mathbf{B}) - \frac{\lambda}{2N} \rho(\mathbf{W}, \mathbf{B})$$

# Regularization

- Recall

$$\min_{\mathbf{W}, \mathbf{B}} R_{\mathcal{D}}(\mathbf{W}, \mathbf{B}) - \frac{\lambda}{2N} \rho(\mathbf{W}, \mathbf{B})$$

- L2 regularization:  $\rho(\mathbf{W}, \mathbf{B}) = \|\mathbf{W}\|_2^2 = \sum_{l,i,j} (w_{ij}^l)^2$

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}} - \alpha \frac{\lambda}{N} \mathbf{W} = \left(1 - \frac{\alpha \lambda}{N}\right) \mathbf{W} - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}}$$

# Regularization

- Recall

$$\min_{\mathbf{W}, \mathbf{B}} R_{\mathcal{D}}(\mathbf{W}, \mathbf{B}) - \frac{\lambda}{2N} \rho(\mathbf{W}, \mathbf{B})$$

- L2** regularization:  $\rho(\mathbf{W}, \mathbf{B}) = \|\mathbf{W}\|_2^2 = \sum_{l,i,j} (w_{ij}^l)^2$

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}} - \alpha \frac{\lambda}{N} \mathbf{W} = \left(1 - \frac{\alpha \lambda}{N}\right) \mathbf{W} - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}}$$

- L1** regularization:  $\rho(\mathbf{W}, \mathbf{B}) = \|\mathbf{W}\|_1 = \sum_{l,i,j} |w_{ij}^l|$

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \frac{\lambda}{N} \text{sign}(\mathbf{W}) - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}}$$

# Regularization

- Recall

$$\min_{\mathbf{W}, \mathbf{B}} R_{\mathcal{D}}(\mathbf{W}, \mathbf{B}) - \frac{\lambda}{2N} \rho(\mathbf{W}, \mathbf{B})$$

- L2** regularization:  $\rho(\mathbf{W}, \mathbf{B}) = \|\mathbf{W}\|_2^2 = \sum_{l,i,j} (w_{ij}^l)^2$

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}} - \alpha \frac{\lambda}{N} \mathbf{W} = \left(1 - \frac{\alpha \lambda}{N}\right) \mathbf{W} - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}}$$

- L1** regularization:  $\rho(\mathbf{W}, \mathbf{B}) = \|\mathbf{W}\|_1 = \sum_{l,i,j} |w_{ij}^l|$

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \frac{\lambda}{N} \text{sign}(\mathbf{W}) - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}}$$

- Intuition:** Less sensitive to noise

# Regularization

- Recall

$$\min_{\mathbf{W}, \mathbf{B}} R_{\mathcal{D}}(\mathbf{W}, \mathbf{B}) - \frac{\lambda}{2N} \rho(\mathbf{W}, \mathbf{B})$$

- L2** regularization:  $\rho(\mathbf{W}, \mathbf{B}) = \|\mathbf{W}\|_2^2 = \sum_{l,i,j} (w_{ij}^l)^2$

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}} - \alpha \frac{\lambda}{N} \mathbf{W} = \left(1 - \frac{\alpha \lambda}{N}\right) \mathbf{W} - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}}$$

- L1** regularization:  $\rho(\mathbf{W}, \mathbf{B}) = \|\mathbf{W}\|_1 = \sum_{l,i,j} |w_{ij}^l|$

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \frac{\lambda}{N} \text{sign}(\mathbf{W}) - \alpha \nabla_{\mathbf{W}} R_{\mathcal{D}}$$

- Intuition:** Less sensitive to noise
- Bias  $\mathbf{B}$  generally not regularized. Why?

# Dropout

- **Principle:**
  - Training: at each gradient step, remove temporally half of the nodes of the ANN at random and use the remaining for training
  - Inference: divide weights from hidden nodes by 2
- **Intuition:** more robust learning



# Data Augmentation Techniques (Not Only for ANNs)

- More data  $\Rightarrow$  better trained model, but data costly to collect
- **Idea:** artificially expand dataset at hand
- Input noise
  - generate new sample  $(\mathbf{x} + \varepsilon, \mathbf{y})$  where  $(\mathbf{x}, \mathbf{y}) \in \mathcal{D}$  and  $\varepsilon$  some (Gaussian) random noise
- Problem (roughly) invariant for some transform  $f : \mathcal{X} \text{ to } \mathcal{X}$ 
  - generate new sample  $(f(\mathbf{x}), \mathbf{y})$  where  $(\mathbf{x}, \mathbf{y}) \in \mathcal{D}$
  - for images,  $f =$  (small) translation, rotation, zoom in/out...